

# Supplementary Material: What Does a Computer-Use Agent Reliability Score Actually Measure?

Anonymous Author(s)

**Table 1: Per-task values on the four-task common support.  $\Delta$  is Flash minus GPT-5.5. Task IDs as registered.**

| Task      | $S^0_G$ | $S^0_F$ | STSG  | STSF  | $\Delta S^0$ | $\Delta STS$ |
|-----------|---------|---------|-------|-------|--------------|--------------|
| cf-f010   | 83      | 100     | 0.000 | 0.000 | +17          | 0.000        |
| pref-f014 | 0       | 84      | 0.000 | 0.000 | +84          | 0.000        |
| retr-f002 | 75      | 100     | 0.500 | 0.500 | +25          | 0.000        |
| retr-f009 | 40      | 100     | 0.500 | 0.333 | +60          | -0.167       |

**Table 2: Flash-minus-GPT  $\Delta STS$  on the four-task support.**

| Config  | STS GPT | STS Flash | $\Delta STS$ | Outcome       |
|---------|---------|-----------|--------------|---------------|
| FROZEN  | 0.2500  | 0.2083    | -0.0417      | baseline      |
| R-AGG   | 0.3333  | 0.3125    | -0.0208      | narrowed mean |
| R-SCOPE | 0.2083  | 0.1875    | -0.0208      | narrowed mean |
| R-CMP   | 0.2500  | 0.2083    | -0.0417      | unchanged     |
| R-CHAN  | 0.0417  | 0.2083    | +0.1667      | inverted      |
| ALL     | 0.0417  | 0.2708    | +0.2292      | inverted      |

## Abstract

Optional supplementary material for the AAMAS 2027 main-track submission. Per the call, this file does not count against the eight-page main-text limit, reviewers need not consult it, and it is not an extended or corrected version of the paper. Protocol tables only. No new experiment, no new estimand, no pooled  $n$ . The main paper is self-contained.

## ACM Reference Format:

Anonymous Author(s). 2027. Supplementary Material: What Does a Computer-Use Agent Reliability Score Actually Measure?. In *Proc. of the 26th International Conference on Autonomous Agents and Multiagent Systems (AAMAS 2027), Hanoi, Vietnam, 3–7 May 2027*, IFAAMAS, 1 page.

## 1 Study 2 per-task common support

Numbers already used in the main text; this table is the per-task expansion of the exploratory four-task slice ( $n_{\text{eff}}=1$ ). It is not a ranking and is not the confirmatory selection test.

## 2 Study 3 ordering under repairs

The main-text claims use R-AGG (+8/-18) and ALL (+10/-20) only. The rows below are the same locked counterfactuals on the four-task support; FROZEN is the baseline, not a finding. R-CHAN inversion is an existence demonstration ( $n_{\text{eff}}=1$ ), not a recommended scoring rule.

R-AGG’s mean gap shrinks while per-task disagreement grows: nonzero  $\Delta STS$  tasks go  $1 \rightarrow 2$ , and the two moves cancel (counterfactual-f010 +0.250, retrieval-f009 -0.333).

## 3 Reconstruction (G2)

Locked rule  $R$ , G1 groundedness 30/30, G2 calibration 9/30. MATCH 20  $\rightarrow$  13, RECALL.MISS 39  $\rightarrow$  46, ABSENT 61 unmoved.  $R$  is not revised. The sealed corpus was not opened. G2 is not a result in the main paper.

## 4 Comparative cohort gate

Pre-specified threshold:  $n \geq 20$  independent task clusters, same construct, frozen evaluator, decided before the first trajectory. Slate: 28 new family tasks (not the C7 support, not preference, not sealed). Probe: inject/guest feasibility only; no agent, no STS. Wave A 9/12, Wave B 7/16,  $n = 16 < 20$ . FAIL. Comparative branch closed. Not a scientific headline.



This work is licensed under a Creative Commons Attribution International 4.0 License.

*Proc. of the 26th International Conference on Autonomous Agents and Multiagent Systems (AAMAS 2027)*, M. Baldoni, F. Fang, W. Yeoh, N. Yorke-Smith (eds.), 3–7 May 2027, Hanoi, Vietnam. © 2027 International Foundation for Autonomous Agents and Multiagent Systems (www.ifaamas.org).

## 5 What this supplement is not

Not a place to pool 24 + 57 + 134. Not a ranking. Not a recommended repair. Not a G2 rescue. Not an extended version of the main paper.