
When the World Changes but the Score Does Not: A Counterfactual Audit of Computer-Use Agents

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Computer-use agents are typically scored by whether an episode finishes and what
2 a rubric judge assigns. That scalar mixes three questions: termination, whether
3 the final answer reflected the world the agent saw, and whether the judge noticed
4 when that world changed. We separate them with paired counterfactual episodes
5 on MyPCBench: a locked intervention rewrites a pre-registered determining set D
6 while instruction, interface, and judge stay fixed. Tracking is read from the final
7 answer against guest gold, never from the score; score sensitivity is exact equality
8 $S_{CF} = S_{Base}$ versus any nonzero delta. Across agents, successful counterfactual
9 tracking sometimes leaves the benchmark score unchanged and sometimes
10 changes it; conversely, unchanged high scores can occur despite incomplete state
11 tracking. On a frozen 5×10 design (50 intended cells; 4 never scheduled; 46 ex-
12 ecuted; 92 legs; 35 execution failures; 24 valid pairs), Claude has 9 valid pairs (7
13 tracking-valid: 6 Type A, 1 score-sensitive, 2 Type B; invariance $6/7 = 0.857$,
14 Clopper-Pearson 95% CI $[0.421, 0.996]$); GPT-5.5 has 8 valid (7 tracking-valid:
15 3 Type A, 4 score-sensitive, 1 Type B; $3/7 = 0.429$, $[0.099, 0.816]$); Qwen3.5-
16 35B-A3B has one valid pair (Type A; $[0.025, 1.000]$). Qwen3.5-9B (ablation) and
17 Qwen3.8-Flash (exploratory) are not pooled into these rates. A stable score does
18 not certify tracking; a moving score does not certify a tracking failure.

19 1 Introduction

20 Leaderboard evaluation of computer-use agents (CUAs) reports a completion score: the episode
21 ended, and a rubric assigned a number [Jang et al., 2026, Xie et al., 2024]. That number is asked to
22 stand in for three facts: whether the harness recorded a terminal DONE; whether the answer matched
23 the environment that was present; and whether the judge’s score moved when that state moved.

24 These facts need not agree. An agent can finish, miss a manipulated field, and still score 100 if the
25 rubric does not pin that field. An agent can finish, report the field correctly on both sides of a paired
26 comparison, and still see the score swing for reasons orthogonal to that field. Inferring tracking or
27 judge sensitivity from the scalar alone is therefore undefined.

28 This paper makes the three facts separately observable. For each frozen task we run a base episode
29 in E^0 and a counterfactual episode in $E^1 = I(E^0)$, where I is a locked SQL (and, when designed,
30 file) patch on a pre-registered determining set D . Prompt, UI, screenshot size, and judge are held
31 fixed. Guest probes confirm that gold moved. We code *completion* from the last trajectory action,
32 *tracking* from the final answer against guest gold, and *score sensitivity* from exact scalar equality.
33 Incomplete runs are execution failures, not tracking misses.

34 **Contributions.** (1) An audit taxonomy of valid pairs: Type A (tracked, score-invariant), score-
35 sensitive (tracked, score moves), and high-score tracking failure (Type B: incomplete tracking with

Table 1: Taxonomy of valid pairs. Rows: agent vs. gold. Columns: exact S . Execution failures are excluded.

	Score invariant	Score moves
Tracks D	Type A	Score-sensitive
Misses D	Type B (high, invariant)	not a headline class here

a high, invariant score). (2) Empirical separation on a frozen ten-task MyPCBench audit across three primary lanes, with a size ablation and an exploratory lane held out of the primary rates. (3) Pre-registered intervention roles predicted what state was manipulated, not how the agent would respond to that manipulation. The claim is methodological and local: we do not estimate a population failure rate and we do not headline a pooled invariance fraction.

2 Setup

Intervention. For task t , E_t^0 is the seeded guest and $E_t^1 = I_t(E_t^0)$ is the same guest after a locked patch on D_t . Some tasks are dual-channel (sqlite and a paired local file must move together); others hold one channel constant as a distractor. The same agent A , instruction T_t , and judge J produce $\tau^0 = A(T_t, E_t^0)$ and $\tau^1 = A(T_t, E_t^1)$. Every intervention in this snapshot has a confirming guest probe.

Reads. $G(\tau)$ is the task-relevant quantities in the final answer. $S(\tau) = J(\tau) \in \{0, \dots, 100\}$. Writer-side track fields, `result.txt`, and the mere presence of a score file are not used for DONE, tracking, or validity. `done(\tau) = 1` iff the last recorded action is literally DONE. A pair is valid iff both legs are done. `track = 1` iff $G(\tau^0)$ matches D_t on E_t^0 and $G(\tau^1)$ matches D_t on E_t^1 ; joint D_t requires every component on both legs. Tracking is never inferred from S . On a tracked pair, invariance is $S(\tau^1) = S(\tau^0)$ with no tolerance.

Classes. These labels are an audit taxonomy for this protocol, not a claim that CUAs have three natural kinds. Type A: `track = 1` and $\Delta S = 0$. Score-sensitive: `track = 1` and $\Delta S \neq 0$. Type B (high-score tracking failure): `track = 0` with a high, invariant score. The high-score cutoff used here is operational: both legs ≥ 80 and $\Delta S = 0$ on the three primary Type B cells (100→100 on the ablation/exploratory Type B cells). A different cutoff would relabel the same trajectories; we report the scores, not a law of nature. Type A is not automatically a rubric bug; Section 3.6 reads per-criterion arrays where they exist.

Universe. Ten tasks from one pre-registered eligibility pool (Table 2). Designed roles were frozen before runs. Lanes: Claude, GPT-5.5, Qwen3.5-35B-A3B (primary); Qwen3.5-9B (ablation); Qwen3.8-Flash (exploratory). Design matrix: 50 intended cells. Never scheduled (not execution failures): Qwen3.5-9B and Qwen3.8-Flash on `preference_inference-f018` and `counterfactual-f004` (4 cells). Executed: 46 cells, 92 legs; 57 DONE, 35 execution failures (15 of 35 on Qwen3.5-35B-A3B, mostly empty tool XML). Valid pairs: 24. Invariance rate = Type A / tracking-valid. Intervals are exact Clopper–Pearson 95% CIs. Lanes are never pooled.

Compute. This paper does not train models and performs no additional agent runs. Claude and GPT-5.5 cells used provider computer-use APIs; Qwen cells used local open-weight inference. Original wall-clock, GPU type, and cluster identifiers were not instrumented for this audit.

3 Results

3.1 Type A: tracked, score unchanged

All five lanes complete `retrieval-f001` (loyalty tier/miles Gold/38,450 → Silver/8,620) and track both legs; scores stay at the base value (100 on four lanes; 80 on Qwen3.5-9B). The rubric does not pin Gold or 38,450.

Table 2: Frozen tasks and pre-registered design role.

Task	Phase	Designed role
retrieval-f001	A	control
aggregation-f003	A	control
preference_inference-f018	A	designed Type B (joint D)
counterfactual-f004	A	score-sensitive contrast
retrieval-f003	B	control
retrieval-f016	B	control
retrieval-f029	B	control (dual-channel)
retrieval-f030	B	designed Type B
aggregation-f018	B	designed Type B
preference_inference-f004	B	score-sensitive contrast

Claude aggregation-f018 was designed as a Type B trap (sqlite charitable \$950 $\rightarrow$ \$0 and home-office 104 $\rightarrow$ 0 days; W-2/1099 and files held). Claude is the only lane that completed both legs. Observed class is Type A: CF zeroes both sqlite fields, holds the untouched file, and flags the mismatch; 100 $\rightarrow$ 100. This is the cleanest positive control.

Other primary Type A pairs: Claude retrieval-f003 65 $\rightarrow$ 65, retrieval-f016 100 $\rightarrow$ 100, retrieval-f029 100 $\rightarrow$ 100, aggregation-f003 80 $\rightarrow$ 80; GPT retrieval-f001, retrieval-f003 100 $\rightarrow$ 100, aggregation-f003 50 $\rightarrow$ 50; Qwen3.5-35B-A3B retrieval-f001 100 $\rightarrow$ 100.

3.2 Score-sensitive: tracked, score moves

preference_inference-f004 flips HangryDash from Cooper’s Seafood House (88) to Backyard Ale House (59), TableFind held. Claude and GPT both report the flip. Scores drop 100 $\rightarrow$ 79 and 100 $\rightarrow$ 58 because the rubric pins Cooper’s as HD top.

GPT also tracks three retrieval interventions whose scores move: retrieval-f016 100 $\rightarrow$ 85; retrieval-f029 33 $\rightarrow$ 100; retrieval-f030 53 $\rightarrow$ 100 (TY2025 1099 held at \$1,200; charitable \$950 $\rightarrow$ \$100). Type A and score-sensitive co-exist for the same agent and protocol.

3.3 Type B: high-score tracking failure

1. Claude retrieval-f030, 100 $\rightarrow$ 100. Base answers filed TY2024 1099-NEC \$1,080, not determining TY2025 \$1,200. CF reports \$1,200 and charitable \$100. The miss is the base-leg year, not CF charitable.

2. Claude preference_inference-f018, 100 $\rightarrow$ 100. Joint D : GME 85 $\rightarrow$ 0 and OddsMarket YES 200/active $\rightarrow$ 0/settled. CF lists GME 0 and still lists GameStop YES as 200 / \$16.00.

3. GPT counterfactual-f004, 87 $\rightarrow$ 87. Guest zeroes the 1099; CF still reports a \$1,200 stipend and \$1,068/year net.

Ablation/exploratory Type B (not primary): Qwen3.5-9B retrieval-f016 base \$8,788.75 vs. gold \$8,213.25 (CF exact), 100 $\rightarrow$ 100; Qwen3.8-Flash retrieval-f029 CF “Total Income \$91,200” vs. Box 1 \$90,000, 100 $\rightarrow$ 100.

3.4 Designed role vs. observed class

Pre-registered intervention roles predicted what state was manipulated, not how the agent would respond to that manipulation. aggregation-f018 (designed Type B) $\rightarrow$ Claude Type A. retrieval-f030 (designed Type B) $\rightarrow$ Claude Type B and GPT score-sensitive. preference_inference-f018 $\rightarrow$ Claude Type B as specified. preference_inference-f004 $\rightarrow$ Claude and GPT sensitive. counterfactual-f004 (designed sensitive) $\rightarrow$ GPT Type B.

3.5 Per-model counts

Qwen3.5-35B-A3B’s 1/1 is execution coverage (nine of ten cells fail), not an invariance finding.

Table 3: Primary tier. Invariance = Type A / tracking-valid. Exact Clopper–Pearson 95% CI. $n \leq 9$.

Model	Valid	Tracking-valid	Type A	Sensitive	Type B	Invariance	95% CI
Claude	9	7	6	1	2	0.857	[0.421, 0.996]
GPT	8	7	3	4	1	0.429	[0.099, 0.816]
Qwen3.5-35B-A3B	1	1	1	0	0	1.000	[0.025, 1.000]

Table 4: Ablation and exploratory. Not pooled with Table 3.

Model	Valid	Track.	A	Sens.	B	Inv.	95% CI
Qwen3.5-9B	3	2	1	1	1	0.500	[0.013, 0.987]
Qwen3.8-Flash	3	2	1	1	1	0.500	[0.013, 0.987]

3.6 Judge vs. agent

On aggregation-f003 (\$4,871.70 $\rightarrow$ \$400), four completing lanes track both totals. Claude and GPT are Type A; 9B and Flash are score-sensitive. GPT stays at [1, 1, 0, 0] both legs: refund-linked criteria pass twice; year-count criteria never fire. 9B and Flash each pick up one extra criterion under CF. Type A is not “the refund criteria were blind”; score-sensitive is not “the agent failed to track.”

4 Discussion

Completion scores measure termination plus whatever the rubric binds. They do not measure whether D was read, and they do not measure whether J is attached to D . Type A shows tracking without score movement; score-sensitive shows both; Type B is a high-score tracking failure under this taxonomy, not a fundamental category of CUA behavior.

Score invariance is not a defect to patch out of every rubric. A channel-invariant item that asks whether the agent used the right app and reported a value can award the same score to two correct worlds (retrieval-f001). The same agent can be Type A on some tasks and score-sensitive on others (GPT on retrieval-f001/retrieval-f003/aggregation-f003 vs. retrieval-f016/retrieval-f029/retrieval-f030/preference_inference-f004), so the protocol does not merely detect a globally blind judge.

Broader impacts. Positive: the audit makes hidden evaluation failure modes observable and encourages treating completion scores as insufficient certificates of state tracking. Negative: the same pairing could be used to probe proprietary graders; we intend the results as a diagnostic of scoring, not as an adversarial exploit of any deployed system.

5 Related work

OSWorld [Xie et al., 2024], WebArena [Zhou et al., 2024], and VisualWebArena [Koh et al., 2024] score agents with completion rubrics. MyPCBench [Jang et al., 2026] seeds a personal desktop that the interventions reuse. OSWorld is Apache-2.0; WebArena is Apache-2.0; VisualWebArena is MIT; MyPCBench is MIT with an OSWorld-derived harness under Apache-2.0. None of these leaderboards pair E^0 with $E^1 = I(E^0)$ to split tracking from S . Shao et al. [2026] audit whether intended capability remains necessary for a score; we audit whether a gold-moving intervention appears in the answer and in S . Turk [2026] find that coverage-style scores can hide update failures under controlled mutations; we observe the same pattern for GUI CUAs. Bellibatlu et al. [2026] measure judge instability under prompt paraphrase; we measure sensitivity to a verified world-state change. The pairing follows interventionist language [Pearl, 2009]; we do not identify model internals.

6 Limitations

Ten tasks, one environment, five lanes, $n \leq 9$ valid pairs per primary model. CIs bound binomial uncertainty on this sample, not a deployment rate. Qwen3.5-35B-A3B contributes one semantic




Figure 1: Primary valid pairs: S^0 vs. S^1 . Diagonal: score invariance. Off-diagonal tracked pairs are score-sensitive; high, invariant misses of D are Type B.

144 pair. All 24 valid pairs were hand-coded from final-answer text against guest gold. Score deltas are
 145 attributed to rubric items only where per-criterion arrays exist. Interventions test controlled state
 146 changes, not arbitrary real-world shifts.

147 7 Conclusion

148 On this frozen audit, task completion, counterfactual tracking, and benchmark score sensitivity come
 149 apart. A stable score does not certify tracking; a moving score does not certify a tracking failure.
 150 Paired gold-moving checks make that split observable; CUA leaderboards currently do not run them.

151 References

- 152 Rohith Reddy Bellibatlu, Edward Raff, and Wenbin Zhang. JudgeSense: A benchmark for prompt
 153 sensitivity in LLM-as-a-judge systems, 2026. URL <https://arxiv.org/abs/2604.23478>.
- 154 Lawrence Keunho Jang, Andrew Keunwoo Jang, Jing Yu Koh, and Ruslan Salakhutdinov.
 155 MyPCBench: A benchmark for personally intelligent computer-use agents, 2026.
- 156 Jing Yu Koh, Robert Lo, Lawrence Jang, Vishakh Duvvur, Ming Chong Lim, Po-Yu Huang, Graham
 157 Neubig, Shuyan Zhou, Ruslan Salakhutdinov, and Daniel Fried. VisualWebArena: Evaluating

158 multimodal agents on realistic visual web tasks. In *Proceedings of the 62nd Annual Meeting of*
159 *the Association for Computational Linguistics*, 2024.

160 Judea Pearl. *Causality: Models, Reasoning, and Inference*. Cambridge University Press, 2 edition,
161 2009.

162 Jiaqi Shao, Hanck Chen, Wei Zhang, Maxm Pan, and Bing Luo. Do agent benchmarks measure
163 capability? protocol validity in the age of agentic AI, 2026. URL [https://arxiv.org/abs/](https://arxiv.org/abs/2607.22368)
164 2607.22368.

165 Matt Turk. Counterfactual evaluation reveals hidden capability profiles in clinical LLMs and agents,
166 2026. URL <https://arxiv.org/abs/2605.30590>.

167 Tianbao Xie, Danyang Zhang, Jixuan Chen, Xiaochuan Li, Siheng Zhao, Ruisheng Cao, Toh Jing
168 Hua, Zhoujun Cheng, Daoguang Shin, Fangyu Lei, Yitao Liu, Yiheng Xu, Shuyan Zhou, Silvio
169 Savarese, Caiming Xiong, Victor Zhong, and Tao Yu. OSWorld: Benchmarking multimodal
170 agents for open-ended tasks in real computer environments. In *Advances in Neural Information*
171 *Processing Systems*, 2024.

172 Shuyan Zhou, Frank F. Xu, Hao Zhu, Xuhui Zhou, Robert Lo, Abishek Sridhar, Xianyi Cheng,
173 Tianyue Ou, Yonatan Bisk, Daniel Fried, Uri Alon, and Graham Neubig. WebArena: A realis-
174 tic web environment for building autonomous agents. In *International Conference on Learning*
175 *Representations*, 2024.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: Abstract and Section 1 state a local evaluation claim (separability of completion, tracking, and score sensitivity) matching the frozen-audit results in Section 3, without population-rate claims.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Section 6 discusses the small frozen task set, single environment, $n \leq 9$ per primary model, wide Clopper–Pearson intervals, one semantic pair for Qwen3.5-35B-A3B, and hand-coded tracking.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper contains no theorems or formal proofs; definitions of tracking and invariance are operational coding rules in Section 2.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Section 2 specifies DONE, gold (guest probes), tracking from final answers, exact score equality, the ten-task universe, and validity exclusions; classification criteria are stated per pair in Section 3.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

- (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [No]

Justification: This submission does not include an anonymized code/data release, to avoid deanonymization; the protocol is specified in the paper so classifications can be checked against public benchmark tasks.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Section 2 reports the frozen task list, model lanes, designed roles, coverage (50/46/92/35/24), and scoring/tracking rules needed to interpret Tables 3–4.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Per-model invariance rates use exact Clopper–Pearson 95% intervals (binomial inversion; no $H_0=0.5$ tests), reported in Tables 3–4 and the abstract; $n \leq 9$ is stated.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Section 2 states that the paper does not train models and performs no additional agent runs; Claude and GPT-5.5 used provider computer-use APIs; Qwen used local open-weight inference; original wall-clock, GPU type, and cluster identifiers were not instrumented.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The work is an evaluation-protocol audit on a seeded personal-desktop benchmark with a fictional persona; it does not collect new human-subject data or release models.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Section 4 discusses positive impacts (making hidden evaluation failure modes observable) and negative impacts (the pairing could probe proprietary graders; results are diagnostic, not adversarial).

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The paper does not release pretrained models, generators, or scraped web-scale datasets.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: Section 5 credits the original papers and states licenses: OSWorld Apache-2.0, WebArena Apache-2.0, VisualWebArena MIT, MyPCBench MIT with an OSWorld-derived harness under Apache-2.0. URLs are omitted to preserve anonymity.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [N/A]

Justification: The paper does not release a new dataset, model, or code package with this submission.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: There are no crowdsourcing or human-subject experiments.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: There are no human-subject experiments.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [Yes]

Justification: The evaluated systems and the rubric judge are LLM-based computer-use agents; they are objects of study, not a newly proposed model architecture (Sections 2–3).

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.